

STA141C Final Project

Ryan Cosgrove, Ravinit Chand, Eric Kye, Revanth Rao

2024-06-07

Libraries

```
library(tidyverse)
```

```
## Warning: package 'ggplot2' was built under R version 4.2.3
```

```
## Warning: package 'tidyr' was built under R version 4.2.3
```

```
## Warning: package 'readr' was built under R version 4.2.3
```

```
## Warning: package 'dplyr' was built under R version 4.2.3
```

```
## Warning: package 'stringr' was built under R version 4.2.3
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v ggplot2    3.5.1      v tibble     3.2.1
```

```
## v lubridate  1.9.3      v tidyr      1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Warning: package 'Matrix' was built under R version 4.2.3
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
##
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##     expand, pack, unpack
```

```
##
```

```
## Loaded glmnet 4.1-8
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##     combine
##
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
library(MASS)
```

```
## Warning: package 'MASS' was built under R version 4.2.3
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##     select
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
```

```
library(ROCR)
library(vtable)
```

```
## Loading required package: kableExtra
```

```
## Warning: package 'kableExtra' was built under R version 4.2.3
```

```
##
## Attaching package: 'kableExtra'
##
## The following object is masked from 'package:dplyr':
##
##     group_rows

library(Hmisc)

## Warning: package 'Hmisc' was built under R version 4.2.3

##
## Attaching package: 'Hmisc'
##
## The following objects are masked from 'package:dplyr':
##
##     src, summarize
##
## The following objects are masked from 'package:base':
##
##     format.pval, units
```

Loading Data and Preprocessing

```
weather = read.csv("/Users/eric/Documents/Classroom/23-24/3 Spring/STA141C/Project/weatherAUS.csv")

weather = weather %>% mutate_at(c('WindGustDir', 'WindDir9am', 'WindDir3pm',
                                   'RainToday', 'RainTomorrow'), as.factor)
weather$Year = as.integer(sapply(strsplit(weather[,1], "-"), getElement, 1))
weather_summary = summary(weather)
```

Splitting Data and Removing Variables

```
train_index = (weather$Year < 2013)
test_index = !train_index

train = weather[train_index, ]
test = weather[test_index, ]
```

Remove columns

```
train = train[, c(-1, -2, -8, -10, -11, -24)]
test = test[, c(-1, -2, -8, -10, -11, -24)]
weather = weather[, c(-1, -2, -8, -10, -11, -24)]
weather_plotting = weather[, c(-17, -18)]
```

Table 1: Summary Statistics

Variable	N	Mean	Std. Dev.	Min	Pctl. 25	Pctl. 75	Max
MinTemp	58090	13	6.5	-6.7	8.4	18	31
MaxTemp	58090	24	7	4.1	19	30	48
Rainfall	58090	2.1	7	0	0	0.6	206
Evaporation	58090	5.4	3.7	0	2.8	7.4	81
Sunshine	58090	7.7	3.8	0	5	11	14
WindGustSpeed	58090	41	13	9	31	48	124
WindSpeed9am	58090	15	8.6	0	9	20	67
WindSpeed3pm	58090	20	8.6	0	13	24	76
Humidity9am	58090	66	19	0	55	80	100
Humidity3pm	58090	50	20	0	36	63	100
Pressure9am	58090	1017	6.9	980	1013	1022	1040
Pressure3pm	58090	1015	6.9	977	1010	1020	1039
Cloud9am	58090	4.2	2.8	0	1	7	8
Cloud3pm	58090	4.3	2.7	0	2	7	9
Temp9am	58090	18	6.6	-0.9	13	23	39
Temp3pm	58090	23	6.8	3.7	17	28	46
RainToday	58090						
... No	45323	78%					
... Yes	12767	22%					
RainTomorrow	58090						
... No	45361	78%					
... Yes	12729	22%					


Remove NAs

```
train = na.omit(train)
test = na.omit(test)
weather = na.omit(weather)
RainTom.test <- test$RainTomorrow
```

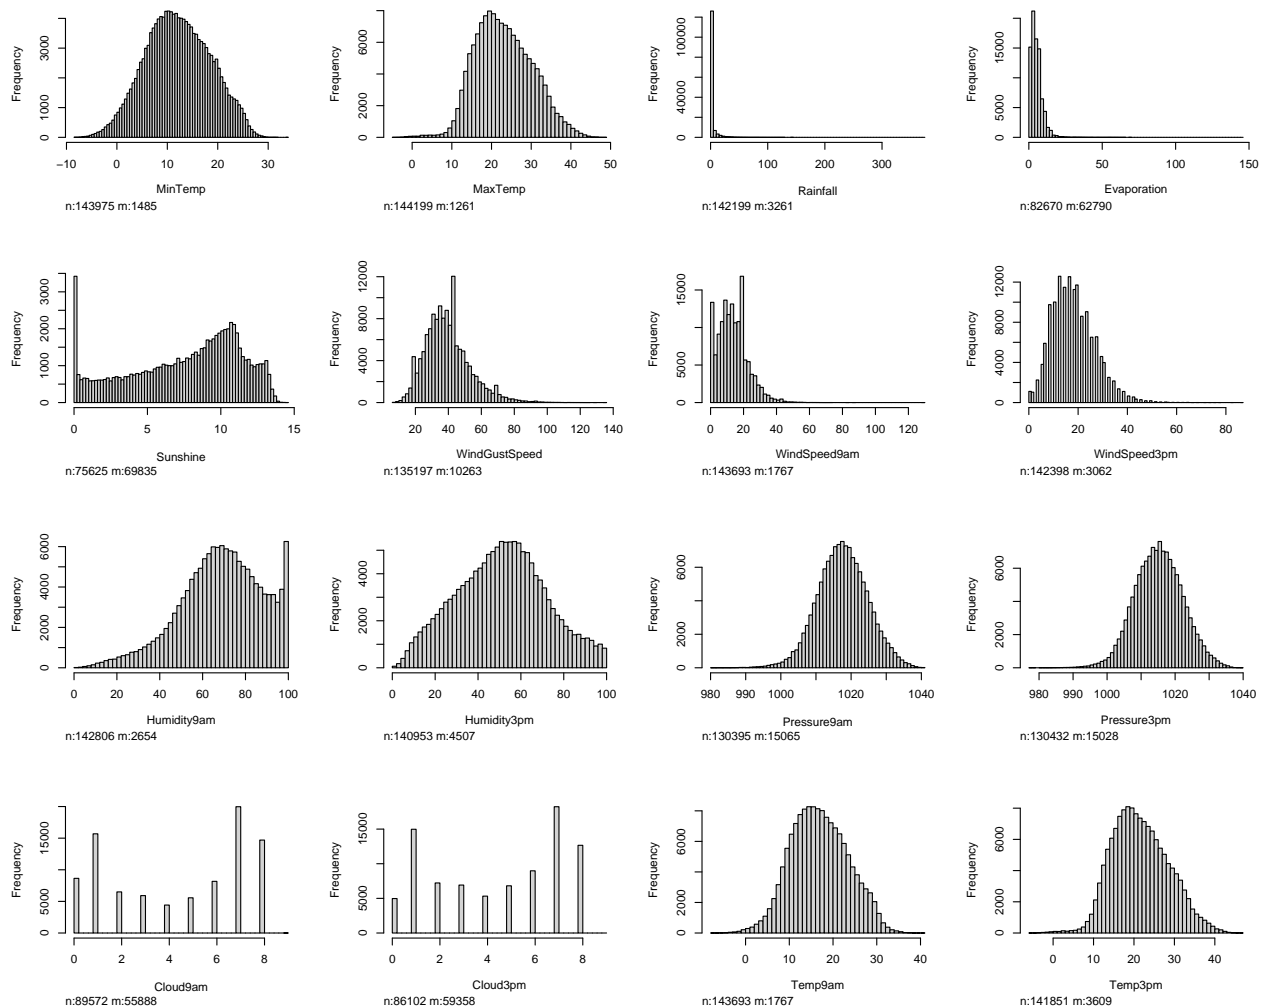
Exploratory Data Analysis

```
st(weather) # Summary Statistics
```

```
corrplot(cor(weather[, c(-17, -18)]), method = "square") # Correlation plot
```



```
# Adjust margins and create histograms of predictor variables
par(mar = c(5, 4, 4, 2) + 0.1)
hist.data.frame(weather_plotting) # Histogram of the Predictor Variables
```



GLM Model

```
glm.fits <- glm(RainTomorrow ~ ., data = train, family = "binomial")
glm.fits
```

```
##
## Call: glm(formula = RainTomorrow ~ ., family = "binomial", data = train)
##
## Coefficients:
## (Intercept)      MinTemp      MaxTemp      Rainfall      Evaporation
##  56.2998671    -0.0478350    -0.0001738     0.0126430    -0.0017503
##    Sunshine WindGustSpeed WindSpeed9am WindSpeed3pm Humidity9am
## -0.1410623     0.0608414    -0.0099919    -0.0282713     0.0020836
## Humidity3pm  Pressure9am  Pressure3pm    Cloud9am    Cloud3pm
##  0.0573718     0.1513636    -0.2137042    -0.0158576     0.1260501
##    Temp9am    Temp3pm RainTodayYes
##  0.0492442     0.0046234     0.4284623
##
## Degrees of Freedom: 31668 Total (i.e. Null); 31651 Residual
```

```
## Null Deviance:      33700
## Residual Deviance: 20990    AIC: 21030

glm.probs <- predict(glm.fits, test, type = "response")
preds= prediction(glm.probs, RainTom.test)
prf = performance(preds, measure = "tpr", x.measure = "fpr")

glm.pred <- rep("No", length(glm.probs))
glm.pred[glm.probs > .5] <- "Yes"
table(glm.pred, RainTom.test)

##           RainTom.test
## glm.pred    No    Yes
##      No  19686  2728
##      Yes   1105  2902

mean(glm.pred == RainTom.test)


## [1] 0.854926

mean(glm.pred != RainTom.test)

## [1] 0.145074
```

GLM plot

```
plot(prf, col = 'red', main = 'ROC Curve for Logistic Regression')
```





LDA Model

```
lda.fit <- lda(RainTomorrow ~ ., data = train)
lda.fit
```

```
## Call:
## lda(RainTomorrow ~ ., data = train)
##
## Prior probabilities of groups:
##      No      Yes
## 0.7758376 0.2241624
##
## Group means:
##      MinTemp MaxTemp Rainfall Evaporation Sunshine WindGustSpeed WindSpeed9am
## No  12.52694 24.25628 1.197285    5.434355 8.557794      38.88225    14.84953
## Yes 14.30008 22.07656 5.815833    4.435526 4.414861      46.34780    16.76884
##      WindSpeed3pm Humidity9am Humidity3pm Pressure9am Pressure3pm Cloud9am
## No      19.23386    64.22145    44.89251    1018.214    1015.675 3.775173
## Yes     21.13706    76.08635    67.62290    1013.794    1011.646 6.022961
##      Cloud3pm Temp9am Temp3pm RainTodayYes
## No  3.824786 17.69324 22.92026    0.1531136
## Yes 6.334272 17.77133 20.16353    0.4699253
##
## Coefficients of linear discriminants:
##              LD1
## MinTemp      -0.043234050
## MaxTemp       0.048836234
## Rainfall      0.013509768
## Evaporation   0.014764696
## Sunshine     -0.132892322
## WindGustSpeed 0.040882713
## WindSpeed9am -0.002372743
## WindSpeed3pm -0.027360965
## Humidity9am  -0.004433778
## Humidity3pm   0.042286151
## Pressure9am   0.095342751
## Pressure3pm  -0.137223770
## Cloud9am     -0.030950057
## Cloud3pm      0.028058901
## Temp9am      -0.001911742
## Temp3pm      -0.007092874
## RainTodayYes  0.427766488
```

```
plot(lda.fit, ylab = "Frequency")
```

```
lda.pred <- predict(lda.fit, test)
```

```
lda.class <- lda.pred$class
table(lda.class, RainTom.test)
```

```
##           RainTom.test
## lda.class   No  Yes
##           No 19592 2657
##           Yes  1199 2973
```

```
mean(lda.class == RainTom.test)
```

```
## [1] 0.8540555
```

```
sum(lda.pred$posterior[, 1] >= .5)
```

```
## [1] 22249
```

```
sum(lda.pred$posterior[, 1] < .5)
```

```
## [1] 4172
```

```
lda.pred$posterior[1:20, 1]
```

```
##           10464           10465           10466           10467           10472           10473           10474
## 0.96799235 0.97727137 0.62769121 0.31341585 0.07218133 0.31421158 0.93979174
##           10478           10479           10480           10481           10488           10490           10492
## 0.08184123 0.31678701 0.73803694 0.78757038 0.89154952 0.19334619 0.21560016
##           10493           10494           10495           10500           10501           10502
## 0.17698203 0.91268072 0.90905031 0.98303784 0.63239428 0.96169231
```

```
lda.class[1:20]
```

```
## [1] No No No Yes Yes Yes No Yes Yes No No No Yes Yes Yes No No No No
## [20] No
## Levels: No Yes
```

```
sum(lda.pred$posterior[, 1] > .9)
```

```
## [1] 15875
```

QDA Model

```
qda.fit <- qda(RainTomorrow ~ ., data = train)
qda.fit
```

```
## Call:
## qda(RainTomorrow ~ ., data = train)
##
## Prior probabilities of groups:
##      No      Yes
## 0.7758376 0.2241624
##
## Group means:
##      MinTemp MaxTemp Rainfall Evaporation Sunshine WindGustSpeed WindSpeed9am
## No  12.52694 24.25628 1.197285    5.434355 8.557794      38.88225    14.84953
## Yes 14.30008 22.07656 5.815833    4.435526 4.414861      46.34780    16.76884
##      WindSpeed3pm Humidity9am Humidity3pm Pressure9am Pressure3pm Cloud9am
## No      19.23386    64.22145    44.89251    1018.214    1015.675 3.775173
## Yes     21.13706    76.08635    67.62290    1013.794    1011.646 6.022961
##      Cloud3pm Temp9am Temp3pm RainTodayYes
## No  3.824786 17.69324 22.92026    0.1531136
## Yes 6.334272 17.77133 20.16353    0.4699253
```

```
qda.class <- predict(qda.fit, test)$class
table(qda.class, RainTom.test)
```

```
##           RainTom.test
## qda.class    No    Yes
##      No  18961  2470
##      Yes   1830  3160
```

```
mean(qda.class == RainTom.test)
```


```
## [1] 0.8372507
```

```
# Recreate x and y after removing NA rows from train and test
x <- model.matrix(RainTomorrow ~ ., rbind(train, test))[, -1]
y <- as.numeric(rbind(train, test)$RainTomorrow) - 1

train_rows <- 1:nrow(train)
test_rows <- (nrow(train) + 1):nrow(x)
```

Lasso Model

```
lasso.fit <- cv.glmnet(x[train_rows, ], y[train_rows], family = "binomial", alpha = 1)
plot(lasso.fit)
```



```
lasso.pred <- predict(lasso.fit, s = "lambda.min", newx = x[test_rows, ], type = "class")
lasso.pred <- ifelse(lasso.pred == "1", "Yes", "No")
table(lasso.pred, RainTom.test)
```

```
##           RainTom.test
## lasso.pred   No  Yes
##           No 19687 2725
##           Yes  1104 2905
```

```
mean(lasso.pred == RainTom.test)
```

```
## [1] 0.8550774
```

```
lasso_coefficients <- predict(lasso.fit, type = "coefficients", s = "lambda.min")
lasso_coefficients[lasso_coefficients != 0]
```


```
## [1] 56.672900914 -0.039524262 0.002861825 0.012524012 -0.002833204
## [6] -0.139631993 0.059845053 -0.010288875 -0.027363712 0.001074917
## [11] 0.057336990 0.141957640 -0.204532156 -0.014934135 0.125236085
## [16] 0.037511549 0.005217262 0.418399522
```

```
length(lasso_coefficients[lasso_coefficients != 0])
```

```
## [1] 18
```

Ridge Model

```
ridge.fit <- cv.glmnet(x[train_rows, ], y[train_rows], family = "binomial", alpha = 0)
plot(ridge.fit)
```



```
ridge.pred <- predict(ridge.fit, s = "lambda.min", newx = x[test_rows, ], type = "class")
table(ridge.pred, RainTom.test)
```

```
##           RainTom.test
## ridge.pred    No  Yes
##           0 19778 2933
##           1  1013 2697
```

```
mean(ridge.pred == RainTom.test)
```

```
## [1] 0
```

```
ridge_coefficients <- predict(ridge.fit, type = "coefficients", s = "lambda.min")
ridge_coefficients[ridge_coefficients != 0]
```

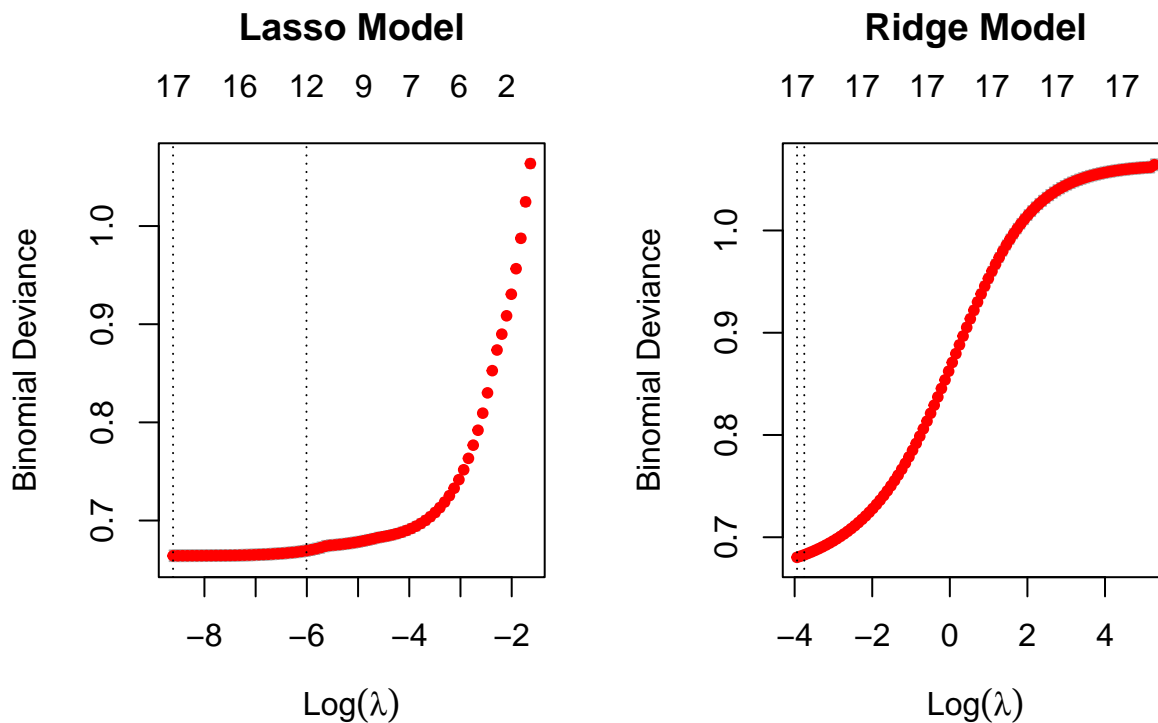
```
## [1] 57.1883694152 -0.0004860718 0.0049484487 0.0129704448 -0.0199356224
## [6] -0.1231279365 0.0383400223 -0.0060877086 -0.0125301733 0.0050073728
## [11] 0.0363669856 -0.0047604579 -0.0565809903 0.0033221501 0.1264372069
## [16] 0.0174424096 -0.0075681767 0.2988933497
```

```
length(ridge_coefficients[ridge_coefficients != 0])
```

```
## [1] 18
```

Plots for Lasso and Ridge Model

```
par(mfrow = c(1, 2), mar = c(5, 4, 6, 2) + 0.1)
plot(lasso.fit, main = "Lasso Model")
plot(ridge.fit, main = "Ridge Model")
```



Random Forest

```
rf.fit <- randomForest(RainTomorrow ~ ., data = train, importance = TRUE)
rf.pred <- predict(rf.fit, newdata = test)
table(rf.pred, RainTom.test)
```

```
##           RainTom.test
## rf.pred    No    Yes
##      No 19808 2766
##      Yes   983 2864
```

```
mean(rf.pred == RainTom.test)
```

```
## [1] 0.8581053
```

Random Forest Importance

```
importance(rf.fit)
```

	No	Yes	MeanDecreaseAccuracy	MeanDecreaseGini
## MinTemp	52.30700	9.2441145	57.15189	521.0837
## MaxTemp	57.32062	0.4328667	59.95045	488.4701
## Rainfall	33.41010	43.0130046	53.63717	570.6358
## Evaporation	57.92732	1.5813540	60.13198	462.0444
## Sunshine	55.94480	78.3050253	91.36357	1286.5645
## WindGustSpeed	76.74449	49.9837049	95.57111	667.5121
## WindSpeed9am	48.44583	-2.3211877	43.71003	382.0238
## WindSpeed3pm	52.58007	2.6907981	50.62741	386.6336
## Humidity9am	55.95183	14.1262398	62.17725	557.1742
## Humidity3pm	79.70070	134.6170325	132.48281	2002.7659
## Pressure9am	58.39126	13.9229813	64.80674	730.8086
## Pressure3pm	74.84276	31.1631977	86.69644	806.5455
## Cloud9am	30.13842	20.2480638	37.39862	340.0611
## Cloud3pm	25.85419	40.8545023	48.25398	621.8783
## Temp9am	60.31213	0.2214884	63.21186	499.1424
## Temp3pm	56.94497	10.4248667	61.91070	511.1755
## RainToday	14.11826	17.4034119	19.78700	175.8386

```
varImpPlot(rf.fit)
```

rf.fit

